

Centro de Investigación en Matemáticas, A.C.

Reconocimiento de Patrones

Jose Miguel Saavedra Aguilar

Examen 2. Ejercicio 2



Para el segundo ejercicio del examen, se implementará un algoritmo Expectation-Maximization (EM) para ajustar una mezcla de k distribuciones gaussianas a un conjunto de datos.

Para este fin, nos basamos en el libro de Bishop, "Pattern Recognition and Machine Learning" del 2009, pues en clase solo estudiamos el algoritmo para variables reales.

```
In [ ]: # Import necessary Libraries for image processing, numerical operations, and statis
from PIL import Image # For image Loading and manipulation
import numpy as np     # For numerical operations and array handling
from scipy.stats import multivariate_normal # For multivariate Gaussian distributi
```

Segun Bishop, para maximizar la verosimilitud de los datos, se deriva la función de log-verosimilitud con respecto a las medias e igualamos a cero, obteniendo la siguiente expresión:

$$\gamma(z_{nk}) = \sum_j \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)$$
$$0 = - \sum_{n=1}^N \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\gamma(z_{nk})} \Sigma_k (x_n - \mu_k).$$

A $\gamma(z_{nk})$ se le conoce como la responsabilidad, o la probabilidad posterior de que el punto x_n pertenezca al componente k de la mezcla, mientras que π_k es el peso, o probabilidad a priori de la mezcla del componente k .

De esta forma, suponiendo que Σ_k es no singular, la expresión anterior se puede reescribir como:

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N \gamma(z_{nk}) x_n,$$

$$N_k = \sum_{n=1}^N \gamma(z_{nk}),$$

podemos interpretar N_k como el número de puntos asignados al componente k de la mezcla.

De forma similar, para la matriz de covarianzas, se sigue un camino análogo derivando con respecto a Σ_k , obteniendo:

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N \gamma(z_{nk}) (x_n - \mu_k)(x_n - \mu_k)^\top.$$

Finalmente, para los pesos de la mezcla, por medio del lagrangiano, se obtiene:

$$\pi_k = \frac{N_k}{N}.$$

En el algoritmo EM, se alternan los pasos de Expectation (E) y Maximization (M) hasta que la convergencia sea alcanzada. En el paso E, se calcula la responsabilidad $\gamma(z_{nk})$ para cada punto x_n y componente k , mientras que en el paso M, se actualizan los parámetros del modelo (medias, covarianzas y pesos) utilizando las responsabilidades calculadas.

A continuación, se presenta la implementación del algoritmo EM para ajustar una mezcla de k distribuciones gaussianas a un conjunto de datos:

```
In [ ]: class GaussianMixtureModel:
    """
    A Gaussian Mixture Model (GMM) implementation using the Expectation-Maximization
    """
    def __init__(self, n_components, max_iters=100, tol=1e-4, random_state=None):
        """
        Initialize the Gaussian Mixture Model.
        Parameters:
        - n_components: Number of Gaussian components in the mixture.
        Optional parameters:
        - max_iters: Maximum number of iterations for the EM algorithm.
        - tol: Convergence tolerance for log-likelihood.
        - random_state: Seed for random number generator for reproducibility.
        """
        # Number of Gaussian components
        self.k = n_components
        # Maximum number of EM iterations
        self.max_iters = max_iters
        # Convergence tolerance for log-likelihood
        self.tol = tol
        # Random seed for reproducibility
        self.random_state = random_state
```

```

def _initialize_parameters(self, X):
    """
    Initialize the parameters of the GMM.
    Randomly selects initial means from the data points and initializes covariances
    """
    # Number of samples and features
    n_samples, n_features = X.shape
    if self.random_state is not None:
        np.random.seed(self.random_state)

    # Randomly choose initial means from data points
    indices = np.random.choice(n_samples, self.k, replace=False)
    self.means_ = X[indices]
    # Initialize covariances to the sample covariance, regularized
    self.covariances_ = np.array([np.cov(X.T) + 1e-6 * np.eye(n_features) for _ in range(self.k)])
    # Initialize equal priors
    self.prior_ = np.full(self.k, 1 / self.k)

def _e_step(self, X):
    """
    Perform the Expectation step of the EM algorithm.
    Computes the posterior probabilities (responsibilities) for each component

    Args:
        X (np.ndarray): Input data of shape (n_samples, n_features).
    """
    n_samples = X.shape[0]
    # Posterior probabilities (responsibilities)
    self.posterior_ = np.zeros((n_samples, self.k))

    for i in range(self.k):
        # Compute probability of each point under component i
        rv = multivariate_normal(mean=self.means_[i], cov=self.covariances_[i])
        self.posterior_[ :, i] = self.prior_[i] * rv.pdf(X)

    # Normalize responsibilities across components
    self.posterior_ /= self.posterior_.sum(axis=1, keepdims=True)

def _m_step(self, X):
    """
    Perform the Maximization step of the EM algorithm.
    Updates the parameters (means, covariances, and priors) based on the current posterior
    Args:
        X (np.ndarray): Input data of shape (n_samples, n_features).
    """
    n_samples, n_features = X.shape
    # Effective number of points assigned to each component
    Nk = self.posterior_.sum(axis=0)
    # Update priors
    self.prior_ = Nk / n_samples
    # Update means
    self.means_ = np.dot(self.posterior_.T, X) / Nk[:, np.newaxis]

    self.covariances_ = []
    for i in range(self.k):
        # Compute weighted covariance for each component

```

```

        diff = X - self.means_[i]
        weighted_diff = self.posterior[:, i][:, np.newaxis] * diff
        cov = weighted_diff.T @ diff / Nk[i]
        cov += 1e-6 * np.eye(n_features) # regularization for numerical stabil
        self.covariances_.append(cov)
    self.covariances_ = np.array(self.covariances_)

def _gaussian_log_likelihood(self, X):
    """
    Compute the log-likelihood of the data given the current model parameters.
    Args:
        X (np.ndarray): Input data of shape (n_samples, n_features).
    """
    n_samples = X.shape[0]
    likelihood = np.zeros((n_samples, self.k))
    for i in range(self.k):
        # Compute Likelihood of each point under each component
        rv = multivariate_normal(mean=self.means_[i], cov=self.covariances_[i])
        likelihood[:, i] = self.prior[i] * rv.pdf(X)
    # Return total Log-Likelihood
    return np.sum(np.log(likelihood.sum(axis=1)))

def fit(self, X):
    """
    Fit the Gaussian Mixture Model to the data using the EM algorithm.

    Args:
        X (np.ndarray): Input data of shape (n_samples, n_features).
    """
    # Initialize parameters
    self._initialize_parameters(X)
    self.log_likelihood_ = -np.inf

    for _ in range(self.max_iters):
        # E-step: update responsibilities
        self._e_step(X)
        # M-step: update parameters
        self._m_step(X)
        # Compute Log-Likelihood
        new_log_likelihood = self._gaussian_log_likelihood(X)
        # Check for convergence
        if abs(new_log_likelihood - self.log_likelihood_) < self.tol:
            break
        self.log_likelihood_ = new_log_likelihood

def predict_proba(self, X):
    """
    Predict the posterior probabilities of each component for the given data.
    Args:
        X (np.ndarray): Input data of shape (n_samples, n_features).
    Returns:
        np.ndarray: Posterior probabilities of shape (n_samples, n_components).
    """
    n_samples = X.shape[0]
    probs = np.zeros((n_samples, self.k))
    for i in range(self.k):

```

```

        # Compute probability of each point under each component
        rv = multivariate_normal(mean=self.means_[i], cov=self.covariances_[i])
        probs[:, i] = self.prior_[i] * rv.pdf(X)
    # Normalize to get probabilities
    probs /= probs.sum(axis=1, keepdims=True)
    return probs

def predict(self, X):
    """
    Predict the component labels for the given data based on the highest poster
    Args:
        X (np.ndarray): Input data of shape (n_samples, n_features).
    Returns:
        np.ndarray: Predicted component labels of shape (n_samples,).
    """
    # Assign each point to the component with highest probability
    return np.argmax(self.predict_proba(X), axis=1)

def fit_predict(self, X):
    """
    Fit the model to the data and predict the component labels.
    Args:
        X (np.ndarray): Input data of shape (n_samples, n_features).
    Returns:
        np.ndarray: Predicted component labels of shape (n_samples,).
    """
    self.fit(X)
    return self.predict(X)

def get_parameters(self):
    """
    Get the parameters of the GMM.
    Returns:
        dict: Dictionary containing means, covariances, and priors.
    """
    return {
        'means': self.means_,
        'covariances': self.covariances_,
        'priors': self.prior_
    }

```

Definimos funciones auxiliares para convertir la imagen a datos, y la función inversa:

```

In [ ]: # Convert an image to a feature array for clustering
# Each pixel is represented as a 3-dimensional vector (RGB)
def image_to_features(image_path):
    img = Image.open(image_path).convert('RGB') # Load image and ensure RGB format
    img_np = np.array(img) # Convert image to numpy array
    n, m, c = img_np.shape # Get image dimensions
    features = img_np.reshape(-1, 3) # Flatten image to (num_pixels, 3)
    return features, (n, m)

# Convert clustered features back to image format
def features_to_image(features, shape):

```

```
n, m = shape # Unpack original image shape
return features.reshape(n, m, 3).astype(np.uint8) # Reshape and convert to uint8
```

Para los 4 valores de k , (2, 3, 5 y 10), se ejecuta el algoritmo de E-M, o de mezcla de Gaussianas (GMM) para colorear la imagen:

```
In [ ]: # List of different numbers of gaussians to try
K = [2, 3, 5, 10]

for k in K:
    # Load and transform the image into features (pixels as RGB vectors)
    features, shape = image_to_features("Figures/foto.jpg")

    # Fit the Gaussian Mixture Model to the pixel features
    em = GaussianMixtureModel(n_components=k, max_iters=30, tol=1e-4, random_state=0)
    labels = em.fit_predict(features) # Cluster assignment for each pixel

    parameters = em.get_parameters() # Get Learned means

    output = np.zeros_like(features) # Prepare array for reconstructed image

    # Assign each pixel the mean color of its cluster
    for i in range(shape[0] * shape[1]):
        output[i] = parameters['means'][labels[i]]

    # Inverse transform: reshape features back to image and save result
    reconstructed_img = features_to_image(output, shape)
    Image.fromarray(reconstructed_img).save(f"gmm_{k}.png")
```

Las imágenes resultantes se encuentran guardadas con el nombre 'gmm_x.png', donde x es el número de gaussianas utilizadas.